

Top 5 European Leagues: Player Performance Analysis (2025-2026)

Dataset: [Football Players Stats \(2025-2026\)](#), sourced from FBref. Covers the Premier League, La Liga, Bundesliga, Serie A, and Ligue 1

How to use this notebook:

1. Download `players_data-2025_2026.csv` from the Kaggle link above and place it in the same folder as this notebook
2. Run all cells. If the real file isn't found, the notebook falls back to a small synthetic sample with the same real column names (Gls, Ast, Min, 90s, etc.), so everything runs end to end before you plug in the real data.

Note on columns: this dataset uses standard FBref abbreviations — `Gls` = Goals, `Ast` = Assists, `Min` = Minutes played, `90s` = minutes / 90 (already pre-calculated), `Pos` = Position, `Squad` = Club, `Comp` = League/Competition.

1. Imports and setup

```
In [49]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import os

sns.set_theme(style="whitegrid")

csv_file = "players_data-2025_2026.csv"
```

2. Load the data

Loads the real Kaggle CSV if it's present. Otherwise, falls back to a small synthetic sample using the **same real column names** as the actual FBref dataset, so the rest of the notebook is fully runnable without the real file, and the code won't need renaming once you plug in the real data.

```
In [50]: def load_data(path: str) -> pd.DataFrame:
    if not os.path.exists(path):
        raise FileNotFoundError(
            f"'{path}' not found. Download the real CSV from Kaggle and place it here"
        )
    print(f"Loading dataset from {path}")
    return pd.read_csv(path)
```

```
df = load_data(csv_file)
df.head()
```

Loading dataset from players_data-2025_2026.csv

```
Out[50]:
```

	Rk	Player	Nation	Pos	Squad	Comp	Age	Born	MP	Starts	...
0	1	Brenden Aaronson	us USA	MF,FW	Leeds United	eng Premier League	24.0	2000.0	37	30	...
1	2	Jerome Abbey	eng ENG	MF	Wolves	eng Premier League	15.0	2009.0	1	0	...
2	3	Zach Abbott	eng ENG	DF	Nottingham Forest	eng Premier League	19.0	2006.0	3	2	...
3	4	Jones El-Abdellaoui	ma MAR	MF,FW	Celta Vigo	es La Liga	19.0	2006.0	22	4	...
4	5	Himad Abdelli	dz ALG	MF	Marseille	fr Ligue 1	25.0	1999.0	8	1	...

5 rows × 102 columns

3. Original look at the data

```
In [51]: df.shape
```

```
Out[51]: (2839, 102)
```

```
In [52]: df.info()
```

```
<class 'pandas.DataFrame'>
RangeIndex: 2839 entries, 0 to 2838
Columns: 102 entries, Rk to OG
dtypes: float64(45), int64(40), str(17)
memory usage: 2.5 MB
```

```
In [53]: df.describe()
```

Out [53]:

	Rk	Age	Born	MP	Starts	
count	2839.000000	2837.000000	2837.000000	2839.000000	2839.000000	2839.000
mean	1420.000000	24.855834	1999.802256	19.138781	13.576611	1218.580
std	819.693032	4.623789	4.633160	11.520385	11.225366	963.218
min	1.000000	15.000000	1983.000000	1.000000	0.000000	1.000
25%	710.500000	21.000000	1997.000000	9.000000	3.000000	320.500
50%	1420.000000	25.000000	2000.000000	20.000000	12.000000	1074.000
75%	2129.500000	28.000000	2003.000000	29.000000	23.000000	1972.500
max	2839.000000	41.000000	2010.000000	38.000000	38.000000	3420.000

8 rows × 85 columns

In [54]:

```
# Let's check the missing values for each column
df.isna().sum().sort_values(ascending=False).head(10)
```

Out [54]:

```
CS%                2651
Save%              2649
SoTA               2645
Comp_stats_keeper  2645
Born_stats_keeper  2645
MP_stats_keeper    2645
Starts_stats_keeper 2645
Min_stats_keeper   2645
90s_stats_keeper   2645
GA                 2645
dtype: int64
```

4. Cleaning the data

- Let's drop rows for players with 0 minutes played as they would break per-90 stats.
- Let's also only keep the columns relevant to the analysis in this notebook, so it stays readable (even though the real dataset has 53 columns total).

In [55]:

```
before = len(df)
df = df[df["Min"] > 0].copy()
after = len(df)
print(f"Dropped {before - after} rows with 0 minutes played.")

keep_cols = ["Player", "Nation", "Pos", "Squad", "Comp", "Age",
             "MP", "Min", "90s", "Gls", "Ast", "G+A", "CrdY", "CrdR"]
keep_cols = [c for c in keep_cols if c in df.columns]
df = df[keep_cols].copy()
df.head()
```

Dropped 0 rows with 0 minutes played.

Out [55]:

	Player	Nation	Pos	Squad	Comp	Age	MP	Min	90s	Gls	Ast	G+
0	Brenden Aaronson	us USA	MF,FW	Leeds United	eng Premier League	24.0	37	2449	27.2	4	5	
1	Jerome Abbey	eng ENG	MF	Wolves	eng Premier League	15.0	1	17	0.2	0	0	
2	Zach Abbott	eng ENG	DF	Nottingham Forest	eng Premier League	19.0	3	164	1.8	0	0	
3	Jones El-Abdellaoui	ma MAR	MF,FW	Celta Vigo	es La Liga	19.0	22	681	7.6	2	0	
4	Himad Abdelli	dz ALG	MF	Marseille	fr Ligue 1	25.0	8	138	1.5	0	0	

5. Feature engineering

Using the dataset's own pre-calculated `90s` column (minutes / 90) to derive **goals per 90** and **assists per 90** — the standard way to compare players who played very different amounts of time.

In [56]:

```
df["Gls_per90"] = df["Gls"] / df["90s"]
df["Ast_per90"] = df["Ast"] / df["90s"]
df[["Player", "Squad", "Gls", "90s", "Gls_per90", "Ast_per90"]].head()
```

Out [56]:

	Player	Squad	Gls	90s	Gls_per90	Ast_per90
0	Brenden Aaronson	Leeds United	4	27.2	0.147059	0.183824
1	Jerome Abbey	Wolves	0	0.2	0.000000	0.000000
2	Zach Abbott	Nottingham Forest	0	1.8	0.000000	0.000000
3	Jones El-Abdellaoui	Celta Vigo	2	7.6	0.263158	0.000000
4	Himad Abdelli	Marseille	0	1.5	0.000000	0.000000


6. Top goal scorers

In [57]:

```
top_n = 10
top = df.nlargest(top_n, "Gls")

plt.figure(figsize=(8, 5))
sns.barplot(data=top, x="Gls", y="Player", hue="Squad", dodge=False)
plt.title(f"Top {top_n} Goal Scorers, 2025-2026")
plt.xlabel("Goals")
plt.ylabel("")
plt.legend(title="Club", bbox_to_anchor=(1.02, 1), loc="upper left")
```

```
plt.tight_layout()
plt.savefig(os.path.join(OUTPUT_DIR, "top_scorers.png"), dpi=150)
plt.show()
```




7. Goal contributions by position

Goals + Assists per 90 minutes, grouped by position — checking whether the data shows the expected pattern (forwards contributing most, defenders least).

```
In [58]: df["GA_per90"] = df["Gls_per90"] + df["Ast_per90"]


plt.figure(figsize=(7, 5))
sns.boxplot(data=df, x="Pos", y="GA_per90")
plt.title("Goal Contributions (G+A) per 90 Minutes, by Position")
plt.xlabel("")
plt.ylabel("G+A per 90")
plt.tight_layout()
plt.savefig(os.path.join(OUTPUT_DIR, "contributions_by_position.png"), dpi=150)
plt.show()
```



8. Age vs. output

Let's check for a visible relationship between a player's age and their per-90 output?
Common football-analytics wisdom says output tends to peak in the mid-to-late 20s.

```
In [59]: plt.figure(figsize=(7, 6))
sns.scatterplot(data=df, x="Age", y="GA_per90", hue="Pos", alpha=0.7)
sns.regplot(data=df, x="Age", y="GA_per90", scatter=False, color="black", li
plt.title("Age vs Goal Contributions per 90")
plt.xlabel("Age")
plt.ylabel("G+A per 90")
plt.legend(bbox_to_anchor=(1.02, 1), loc="upper left")
plt.tight_layout()
plt.savefig(os.path.join(OUTPUT_DIR, "age_vs_output.png"), dpi=150)
plt.show()
```



9. Findings

```
In [60]: top_scorer = df.loc[df["Gls"].idxmax()]
best_per90 = df[df["90s"] >= 5].loc[df[df["90s"] >= 5]["GA_per90"].idxmax()]

print(f"Top scorer: {top_scorer['Player']} ({top_scorer['Squad']}) "
      f"with {top_scorer['Gls']} goals.")

print(f"Best G+A per 90 (min. 5 full matches played): {best_per90['Player']} "
      f"({best_per90['Squad']}), {best_per90['GA_per90']:.2f} goal contribut

print("(The 5-match minimum avoids one-off hot streaks from players with ver
      "playing time dominating the per-90 ranking.)")
```

Top scorer: Harry Kane (Bayern Munich) with 36 goals.

Best G+A per 90 (min. 5 full matches played): Harry Kane (Bayern Munich), 1.55 goal contributions per 90.

(The 5-match minimum avoids one-off hot streaks from players with very little playing time dominating the per-90 ranking.)